

# Deep Learning

## Deep Learning for Text

Chapter 11.1-11.3

Vera Hollink  
2025

# Natural Language Processing (NLP)


- Natural language = human language
  - Vocabulary changes
  - Grammar not well-defined



- NLP
  - do something useful with natural language
  - $\neq$  understanding

# Some NLP Tasks

- Text classification



- Content filtering




- Sentiment analysis



- Translation



- Summarization



# Text preprocessing

# Text preprocessing

The quick brown fox jumps over the lazy dog.

## 1. Normalization/standardization

the quick brown fox jumps over the lazy dog

## 2. Tokenization

"the" "quick" "brown" "fox" "jumps" "over" "the"  
"lazy" "dog"

## 3. Indexing

17 321 490 21 339 3021 17 591 111

## 4. Encoding

[[0,0,0,0,1,0], [[0,1,0,0,0,0], [[1,0,0,0,0,0], ...]

# Normalization/standardization

- Lowercase

The → the

- Remove punctuation

. ? " ...

- Convert special characters

résumé → resume

- Stemming

foxes => fox

approximation → approximat

Disadvantage: information is lost

Advantage: less training data needed

# Tokenization

the quick brown fox jumps over the lazy dog

- Words

"the" "quick" "brown" "fox" "jumps" "over"  
"the" "lazy" "dog"

- N-grams: sequences of N words

2-grams:

"the quick" "quick brown" "brown fox" "fox  
jumps" "jumps over" "over the" "the lazy" "lazy  
dog"

3-grams:

"the quick brown" "quick brown fox" "brown fox  
jumps" "fox jumps over" "jumps over the" "over  
the lazy" "the lazy dog"

- Characters

"t" "h" "e" "q" "u" "i" "c" "k" "b" ...

# Indexing

- Assign number to each token

the → 17

quick → 321

- Use only N most frequent tokens (e.g. 10000)

other tokens get index 1

- Create vectors

[17 321 490 21 339 3021 17 591 111]

- If fixed length needed, pad with 0's:

Length 12:

[17 321 490 21 339 3021 17 591 111 0 0 0]



# Text preprocessing in Keras

Preprocessing module

```
from tensorflow.keras.layers import TextVectorization
dataset = ['The brown dog jumps.', 'Dog jumps over the fox.']
text_vectorization = TextVectorization(output_mode='int')
text_vectorization.adapt(dataset)
text = 'The quick brown dog jumps.'
encoded_text = text_vectorization(text)
print(encoded_text)
```

Default: lowercase,  
remove punctuation,  
split on whitespace  
There are many  
alternative options

Create  
vocabulary: give  
each token a  
number

Apply to  
new data

1 for unknown  
word

```
tf.Tensor([2 1 7 4 3], shape=(5,), dtype=int64)
```

# Text preprocessing

The quick brown fox jumps over the lazy dog.

## 1. Normalization/standardization

the quick brown fox jumps over the lazy dog

## 2. Tokenization

"the" "quick" "brown" "fox" "jumps" "over" "the"  
"lazy" "dog"

## 3. Indexing

17 321 490 21 339 3021 17 591 111

## 4. Encoding

[[0,0,0,0,1,0], [[0,1,0,0,0,0], [[1,0,0,0,0,0], ...]

# Bag-of-word representations

## 1. Multi-hot encoding

Doc id	dog	fox	jump	over	...
1	1	0	1	0	...
2	1	1	0	0	...
3	1	0	0	0	...
4	0	1	0	0	...
5	0	1	1	1	...
...	...	...	...	...	...

```
text_vectorization = TextVectorization(max_tokens = 10000,  
    output_mode='multi-hot')
```

# Bag-of-word representations

## 2. Frequency encoding

Doc id	dog	fox	jump	over	...
1	2	0	1	0	...
2	1	2	0	0	...
3	4	0	0	0	...
4	0	1	0	0	...
5	0	6	2	1	...
...	...	...	...	...	...

```
text_vectorization = TextVectorization(max_tokens = 10000,  
                                       output_mode='count')
```

# Bag-of-word representations

## 3. Tf.idf encoding

Value = importance score of the token

Doc id	dog	fox	jump	over	...
1	2.6	0	0.2	0	...
2	0.9	2	0	0	...
3	3.2	0	0	0	...
4	0	0.9	0	0	...
5	0	2.1	1.9	0.1	...
...	...	...	...	...	...

```
text_vectorization = TextVectorization(max_tokens = 10000,  
                                       output_mode='tf_idf')
```

# Bag-of-word representations

## N-gram encoding (multi-hot)

Doc id	dog_jump	jump_over	over_the	...
1	1	0	0	...
2	0	0	0	...
3	0	0	0	...
4	0	0	0	...
5	0	1	1	...
...	...	...	...	...

```
text_vectorization = TextVectorization(n_grams=2,  
                                       max_tokens = 10000, output_mode='multi-hot')
```

count or tf\_idf  
also possible

# Exercise: text preprocessing and encoding

## Two documents

`"Are oranges always orange?"`

`"The cat ate the oranges."`

- Apply stemming and other normalization techniques
- Apply tokenization
- Encode using
  1. Bag-of-words with multi-hot encoding
  2. Bag-of-words with frequency encoding
  3. Bag-of-words with 2-gram multi-hot encoding

# Exercise: text preprocessing and encoding

Two documents

document 1: "Are oranges always orange?"

document 2: "The cat ate the oranges."

- Apply stemming and other normalization techniques
- Apply tokenization

[are, orange, always, orange]

[the, cat, ate, the, orange]



# Exercise: text preprocessing and encoding

Bag-of-words with multi-hot encoding

Doc id	are	orange	always	the	cat	ate
1	1	1	1	0	0	0
2	0	1	0	1	1	1

Bag-of-words with frequency encoding

Doc id	are	orange	always	the	cat	ate
1	1	2	1	0	0	0
2	0	1	0	2	1	1

# Exercise: text preprocessing and encoding

Bag-of-words with 2-gram multi-hot encoding

```
[are_orange, orange_always, always_orange]
```

```
[the_cat, cat_ate, ate_the, the_orange]
```


Doc id	are_orange	orange_always	always_orange	the_cat	cat_ate	ate_the	the_orange
1	1	1	1	0	0	0	0
2	0	0	0	1	1	1	1

# Deep Learning models for text

1. Dense models
2. Recurrent Neural Networks
3. Transformers (next week)

# Deep Learning models for text

1. Dense models
2. Recurrent Neural Networks
3. Transformers (next week)



# Dense model

- Input: document encoded as bag-of-words

Multi-hot: [0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0...]

Frequency: [0, 0, 2, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 3, 0...]

tf\_idf: [0, 0, 2.4, 0, 0, 0, 0, 0, 0, 0, 0.9, 0, 0, 0, 0, 0, 1.3, 0...]

- Dense model:


Number of words  
in the vocabulary

```
model = Sequential()
model.add(Dense(50, input_shape=(n_words,), activation='relu'))
model.add(Dense(1, activation='sigmoid'))
```

- Disadvantage: word order is lost
  - Partly solved with N-grams, but very short sequences

# Deep Learning models for text

1. Dense models
2. Recurrent Neural Networks
3. Transformers (next week)



# RNN

- Sequence processing: (bidirectional) LSTM

```
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = tf.one_hot(inputs, depth=max_tokens)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
```

- One-hot encoding:

- Each word is a vector with exactly one 1

[0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, ...]

- A sample (document) is a 2-dimensional vector

[[0, 0, 0, 0, 0, 0, 1, 0, 0, , ...], [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...],  
[0, 0, 0, 1, ...], [0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, ...], ...]

- A batch has a number of sequences with the same length

- Use cut-off and padding with 0's
- One batch is for instance 256x20000x600 → training is slow

samples x max\_tokens x sentence\_length

# RNNs and Convnets for text

- (bidirectional) LSTM
- (bidirectional) GRU
- 1D Convolutional Network
- Disadvantage:
  - Slow because input is huge  
samples x max\_tokens x sentence\_length



# One-hot embedding vs word embeddings

- One-hot embedding:

Dog

[0, 0, 0, 0, 0, 0, **1**, 0, 0, 0, 0, ...]

Cat

[0, 0, 0, 0, 0, 0, 0, 0, 0, **1**, 0, ...]




- Word embeddings:

Dog

[0.12, 0.30, 0.20, 0.24]


Cat

[0.73, 0.13, 0.40, 0.44]




# Advantages word embeddings


- Smaller representation
  - E.g. 20000 x 600 becomes 256 x 600
- Vector distances can represent meaning
  - Similar words can have similar vectors  
e.g. vector `cactus` closer to vector `aloe` than to vector `cat`
- Meaningful dimensions
  - Gender, singular/plural, ...
  - Emerging, not hard-coded




# Advantages word embeddings



Male-Female



Verb Tense



Country-Capital

## Learning word embeddings

- Transform word indexes to word vector

[17 321 490 21 339 3021 17 591 111 0 0 0] →

[[0.72, 0.34, 0.1, ...], [0.32, 0.70, ...], ...]

- Initialize randomly, learn to structure space through backpropagation (more similar words get closer vectors)

Length of  
word vector

Number of words  
in sequences

Size of the  
vocabiary

```
inputs = keras.Input(shape=(None,), dtype="int64")
embedded = layers.Embedding(input_dim=max_tokens,
                             output_dim=256, input_length=600, mask_zero=True)(inputs)
x = layers.Bidirectional(layers.LSTM(32))(embedded)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)
model = keras.Model(inputs, outputs)
```

Make sure padding  
(trailing 0's) is not  
used in training.

# Using pretrained word embeddings


- Meaning/embedding not specific for data set or task
- Use existing embedding
  1. Get embedding matrix:  
pretrained vector for each word
  2. Initialize embedding with matrix
  3. Fix embedding (layer not trainable)

```
embedding_layer = layers.Embedding(max_tokens, embedding_dim,  
    embeddings_initializer=keras.initializers.Constant(embedding_m  
    atrix), trainable=False, mask_zero=True  
)
```

# Popular pretrained word embeddings

## word2vec:

- made by Google
- based on news data
- vector length 300



## GloVe:

- made by Stanford,
- based on various sources (various versions available), such as Wikipedia and WWW crawl and Twitter,
- vector length 25-300



# Using pretrained word embeddings

- Advantages?
  - Less training and data needed
  - Based on large corpus
- Disadvantages?
  - Not specialized for your task (e.g. sentiment analysis)
  - Not specialized for your texts (e.g. reviews)

# Deep Learning models for text

1. Dense models
2. Recurrent Neural Networks
3. Transformers (next week)





# Test your understanding!

